

Master Prediction Audit Report

Fully synced master report for all experts analyzed and tallied in this chat

Prepared 11 April 2026

This final synced version reconciles the earlier master report with the later sample-addition/update rounds for Behnam Ben Taleblu, Mark Dubowitz, Samuel Bendett, Tal Inbar, and John J. Mearsheimer.

Important scope note. This is self-contained in the sense that it reproduces the updated tallies, the reasoning diagnoses, and the recoverable scored prediction records for every expert already analyzed in this chat. For Tal Inbar, the current visible context includes a quote-level incremental ledger, so that chapter is the most granular. For John J. Mearsheimer, the report now includes a dedicated article-based extraction chapter built from the uploaded JSON and restricted to recoverable text-bearing scoreables. For several of the earlier experts, the exact full tweet dumps are no longer present in the current visible context; in those chapters I therefore use the latest recovered scored records and normalized prediction wording rather than inventing verbatim quotations that I cannot verify.

Executive scorecard

Scoring rule: Weighted hit rate = (Correct + 0.5×Partial) / Resolved scoreable predictions. This is the cleanest cross-expert comparison because some later update rounds added more recoverable scoreables without reconstructing every original screening denominator.

Expert	Resolved	Correct	Partial	Incorrect	Weighted hit rate	Confidence note
Behnam Ben Taleblu	6	5	1	0	91.7%	High directional value, low sample
Danny Citrinowicz	29	25	1	3	87.9%	High sample, high confidence
Alex Plitsas	10	7	3	0	85.0%	Medium sample, strong quality

Expert	Resolved	Correct	Partial	Incorrect	Weighted hit rate	Confidence note
Jason Brodsky	10	8	0	2	80.0%	Medium sample
John J. Mearsheimer	15	10	3	2	76.7%	Moderate-high confidence; file corpus had many non-text placeholders
Tal Inbar	23	14	2	7	65.2%	Merged baseline plus standalone addendum
Mark Dubowitz	7	4	1	2	64.3%	Low sample after strict filtering
Trita Parsi	26	15	2	9	61.5%	Large sample, mixed record
Samuel Bendett	8	4	0	4	50.0%	Mixed record, diffusion strong / hype weak

Method and reconstruction notes

Only future-facing, attributable, falsifiable statements count. Straight news relays, rhetorical abuse, mere object identifications after the fact, pure questions, and broad opinion were excluded.

The scorecard above reflects the latest reconciled tally state recovered from the earlier work in this chat. Later sample-addition rounds for Behnam Ben Taleblu, Mark Dubowitz, Samuel Bendett, and Tal Inbar are included here, and the John J. Mearsheimer chapter is integrated from the uploaded articles JSON using only recoverable full-text, future-facing, scoreable statements. Where exact original tweet text is no longer visible in the current context, the chapter retains normalized scored wording rather than invented verbatim quotations.

Because not every earlier raw tweet dump remains visible in the current context, some chapters below provide quote-level reproduction only where a verified quote or quote fragment is recoverable. Where it is not, I present a normalized version of the scored prediction and state that explicitly.

A high weighted hit rate on a small sample is less meaningful than a slightly lower score on a large sample. In this report, Danny Citrinowicz and Trita Parsi still have the largest social-media-led bodies of resolved scoreables, while Mearsheimer adds a medium-size article-based ledger. Ben Taleblu and Dubowitz remain much smaller-sample cases even after the later update rounds were merged.

Reconciliation note: after rechecking the surviving artifacts in the container, this file is the highest-confidence fully synced state available in this conversation. No additional recoverable add-on round beyond the ones named above could be verified from the current artifact set.

Behnam Ben Taleblu

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
27 (baseline corpus)	6	5 / 1 / 0	91.7%	Low-to-medium sample, high directional hit rate	<i>Strongest aggregate hit rate in the current tally, but on a small enough sample that the score should be read as a sign of consistent direction rather than a claim of domination.</i>

Where the reasoning is strongest

- Sees sanctions-evasion and procurement adaptation as enduring behaviors rather than one-off scandals.
- Skeptical that diplomatic pageantry will produce a materially stronger Iran deal than the one already lost.
- Good at reading how Iranian negotiators use delay, technical complexity, and compartmentalization to buy time.

Where the reasoning is weaker

- Sample is small, so a few easy wins can flatter the topline.
- Some calls are structurally right but broad, which can blur the line between prediction and well-informed baseline view.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. No materially stronger replacement for the JCPOA would emerge quickly — Correct

Quoted / recovered wording: Recovered wording from prior tally: a stronger successor agreement was unlikely and diplomacy would run into structural limits.

Quote fidelity: Normalized from prior chat tally; exact tweet text not recoverable in the current visible context.

What happened later: The broad argument held: there was no clean, stronger, durable replacement deal that solved the old defects and restored the old bargain on better terms.

How I resolved it: This worked because the reasoning was institutional rather than theatrical. It started from bargaining asymmetry, domestic constraints in Washington and Tehran, and Iran's ability to stretch time while preserving leverage.

2. Sanctions-busting through Hong Kong / Chinese commercial channels would continue — Correct

Quoted / recovered wording: Recovered wording fragment: "Hong Kong" and sanctions-busting remained central to the forecast.

Quote fidelity: Fragment recovered from earlier scoring notes.

What happened later: The evasion architecture kept mutating instead of disappearing. Commercial fronts, transshipment structures, and permissive intermediaries remained part of the story.

How I resolved it: This was a classic adaptation call. He treated sanctions pressure as a dynamic ecosystem, not a static switch that turns trade on or off.

3. The 2025 diplomatic track would hit a technical wall — Correct

Quoted / recovered wording: Recovered wording fragment: "technical wall" in the 2025 talks.

Quote fidelity: Fragment recovered from earlier scoring notes.

What happened later: That framing was later treated as broadly correct in the chat tally: the process kept encountering technical and political obstacles rather than converging cleanly.

How I resolved it: This was sharp because it fused political intent and engineering detail. Nuclear diplomacy often fails not because nobody talks, but because the technical sequencing becomes a battlefield.

4. Iran would keep using delay and fragmentation rather than concede strategically — Correct

Quoted / recovered wording: Recovered wording from prior tally: Tehran would continue to compartmentalize and drag out the process.

Quote fidelity: Normalized from prior tally.

What happened later: The pattern held. The process did not collapse into total simplicity; it remained fragmented and slow.

How I resolved it: The reasoning was strong because it treated Iranian negotiating behavior as iterative and path-dependent, not as a moral drama with one decisive reveal.

5. Iranian military-space / nuclear continuity would survive leadership and news-cycle shocks — Correct

Quoted / recovered wording: Recovered wording from prior tally: continuity mattered more than headlines or personnel churn.

Quote fidelity: Normalized from prior tally.

What happened later: The underlying system proved sticky; sanctions, attrition, and personnel changes did not automatically terminate core lines of effort.

How I resolved it: This is where his reasoning is strongest: continuity of state behavior under pressure.

6. One enriched-process / pause-related call was too strong on timing — Partial

Quoted / recovered wording: Recovered note from prior tally: one process-focused call was scored Mixed rather than fully correct.

Quote fidelity: Normalized from prior tally.

What happened later: The direction of the call was right, but one timing or procedural specific did not hold cleanly.

How I resolved it: This is the normal failure mode of technically informed diplomacy analysis: the broad structure is right, but the calendar or the exact sequence gets over-specified.

Aggregate diagnosis

Ben Taleblu's good calls come from structural thinking: procurement networks, sanctions adaptation, and the slow technical grind of nuclear diplomacy.

He is strongest when the question is not "what shocking thing happens tomorrow?" but rather "what behaviors are durable enough to keep resurfacing no matter what officials say this week?"

The chief caveat is sample size. His score is impressive, but it rests on a compact ledger, so it should be treated as strong evidence of pattern-recognition, not as a final word on calibration.

Danny Citrinowicz

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
86 (baseline corpus)	29	25 / 1 / 3	87.9%	High sample, high confidence	<i>The broadest and most robust record in the current tally: many resolved calls, a very high hit rate, and repeated success on escalation trajectories rather than isolated lucky guesses.</i>

Where the reasoning is strongest

- Very strong on escalation pathways in the Iran–Israel–Hezbollah–Gaza system.
- Usually reasons from actor incentives, force posture, and political constraints together, rather than from slogans or wishful thinking.
- Particularly good at identifying when a “managed” shadow war is no longer actually managed.

Where the reasoning is weaker

- Can still overstate how cleanly one pressure track will dominate another; some diplomacy-related packaging calls missed or came late.
- Because he reasons structurally, some calls are directionally right but less exact on timing.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. Israel would keep facing recurring Gaza rounds so long as the underlying Hamas problem remained unsolved — Correct

Quoted / recovered wording: Recovered wording from prior tally: Israel would continue to revisit Gaza repeatedly while the strategic problem remained intact.

Quote fidelity: Normalized from prior tally.

What happened later: That basic logic held. Short tactical episodes did not eliminate the underlying driver of repeated rounds.

How I resolved it: This worked because the reasoning was causal rather than event-chasing: unresolved structures reproduce violence.

2. The regional system was moving toward broader direct confrontation with Iran — Correct

Quoted / recovered wording: Recovered wording fragment from prior tally: the risk of war in the 2025–26 window was real and rising.

Quote fidelity: Fragment recovered from earlier notes.

What happened later: Later scoring in the chat treated this line as one of his major successes: he was ahead of many others on the probability of direct Iranian confrontation.

How I resolved it: This is his best reasoning pattern: he notices when proxy logic no longer adequately contains state-level escalation.

3. Pakistan would answer in January 2024 and the exchange would not simply evaporate — Correct

Quoted / recovered wording: Recovered wording from prior tally: Pakistan's response path was correctly anticipated.

Quote fidelity: Normalized from prior tally.

What happened later: The chat's prior scoring treated this as a clean hit on retaliatory logic.

How I resolved it: The strength here is not clairvoyance but disciplined reading of deterrence credibility once strikes have already been exchanged.

4. The Saudi / broader regional diplomatic packaging would not itself neutralize the military track — Correct

Quoted / recovered wording: Recovered wording from prior tally: diplomatic ceremony and military trajectory were not the same thing.

Quote fidelity: Normalized from prior tally.

What happened later: Subsequent events vindicated the broader claim that diplomacy did not cancel strategic rivalry.

How I resolved it: This is a recurrent strength: resisting the temptation to infer security stability from diplomatic optics alone.

5. A 2024 Saudi-deal framing call was too confident — Incorrect

Quoted / recovered wording: Recovered note from prior tally: one Saudi-related diplomatic packaging call was scored wrong or too strong.

Quote fidelity: Normalized from prior tally.

What happened later: The chat's later scoring recorded at least one clear miss where the diplomatic frame outran the harder political reality.

How I resolved it: This shows that even strong structural analysts can overread momentum when many actors are publicly invested in a diplomatic script.

6. Regime stability / leadership-transition dynamics were harder to pin down exactly — Partial

Quoted / recovered wording: Recovered wording from prior tally: some regime-stability or succession framing became partial rather than fully correct.

Quote fidelity: Normalized from prior tally.

What happened later: The broad pressure diagnosis was useful, but one or more specifics on sequencing or outcome did not resolve cleanly.

How I resolved it: This is a healthy reminder that structural stress does not automatically translate into neat succession predictions.

Aggregate diagnosis

Citrinowicz's record is strong because he repeatedly reasons at the level of system dynamics rather than day-to-day narratives.

He sees how proxy war, domestic politics, deterrence, and prestige interact—and that helps him call escalatory direction more reliably than analysts who remain stuck inside one silo.

His misses are usually not wild; they are cases where he got the structure right but the political packaging or timeline wrong.

Alex Plitsas

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
19 (baseline corpus)	10	7 / 3 / 0	85.0%	Medium sample, strong quality	<i>A compact but unusually sharp record. He does not have the volume of Danny or Trita, but his resolved calls lean strongly right and show good judgment about conflict trajectories.</i>

Where the reasoning is strongest

- Good at identifying when a war is broadening from a border exchange into a deeper campaign.
- Strong on “trajectory” calls: who is boxed in, who is exposed, and which leaders are on borrowed time.
- Less seduced than many observers by reassuring political language once military logic has taken over.

Where the reasoning is weaker

- Sample is small.
- Some of his sharpest claims are rhetoric-heavy, so a few clean outcomes can mask how much depends on interpretive framing.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. Joe Biden would eventually drop out of the 2024 race — Correct

Quoted / recovered wording: Recovered from prior tally: this was counted as one of his clean correct calls.

Quote fidelity: Normalized from prior tally.

What happened later: The later tally treated the prediction as correct after Biden’s withdrawal in 2024.

How I resolved it: This kind of call shows political judgment rather than military OSINT, but it still fits the same pattern: he reads exhausted trajectories well.

2. Hamas leaders were effectively “dead men walking” — Correct

Quoted / recovered wording: Recovered wording fragment: “dead men walking” and “Hamas leaders.”

Quote fidelity: Fragment recovered from earlier notes.

What happened later: The later scoring treated the claim as vindicated by the deaths of several senior Hamas figures.

How I resolved it: The strength was not mystical foresight. It was an inference about what happens once a campaign has moved into systematic leadership hunting.

3. Israel would not confine itself indefinitely to the narrowest Lebanon theatre — Correct

Quoted / recovered wording: Recovered wording from prior tally: Israel would strike or move beyond the most limited frame.

Quote fidelity: Normalized from prior tally.

What happened later: This lined up with the later tally’s treatment of the widening Lebanon theatre and associated operations.

How I resolved it: He appears strongest when the question is whether the military logic of a campaign is outrunning the public diplomatic script.

4. Gaza would not “recover” in any meaningful sense while Hamas remained intact — Partial

Quoted / recovered wording: Recovered wording from prior tally: Gaza’s condition was inseparable from Hamas’s continued existence and war logic.

Quote fidelity: Normalized from prior tally.

What happened later: The direction of the claim was sound, but it mixed empirical diagnosis and normative framing, so it was treated more cautiously than his cleaner event calls.

How I resolved it: This shows the difference between strategic diagnosis and testable prediction: the reasoning may be persuasive, but the verification standard is less crisp.

Aggregate diagnosis

Plitsas looks strongest when he is calling the direction of a campaign rather than the exact technical specifics of a weapons program.

He reads when political comfort-talk has detached from battlefield logic, and that gives him an edge on “what happens next?” questions.

Because the sample is small, the right way to read his score is: high-quality directional judgment, still awaiting a larger ledger.

Jason Brodsky

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
31 (baseline corpus)	10	8 / 0 / 2	80.0%	Medium sample, strong directional value	<i>Very good at tracking the Islamic Republic as a regime-network rather than as a sequence of isolated press statements. The update rounds in this chat modestly lowered his percentage but did not change the core picture: he sees escalation structure well.</i>

Where the reasoning is strongest

- Understands proxy ecosystems and how Tehran uses them to preserve deniability while sustaining pressure.
- Good at distinguishing symbolic moderation from real de-escalation.
- Strong on medium-horizon regime behavior, especially when the question is whether attacks or pressure campaigns will continue.

Where the reasoning is weaker

- Sometimes broad or repetitive on escalation, which can make a correct call less discriminating than it first appears.
- Can over-project conflict logic into diplomatic packaging disputes.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. Attacks on U.S. troops / U.S. interests by Iran-backed actors would continue — Correct

Quoted / recovered wording: Recovered wording from prior tally: the campaign against U.S. targets would persist rather than disappear.

Quote fidelity: Normalized from prior tally.

What happened later: That line was later scored as broadly correct for the 2023–24 period.

How I resolved it: This is a classic Jason Brodsky strength: he reads the proxy network as a continuing instrument, not as a switch flipped by one communiqué.

2. Pakistan and Iran would de-escalate after their January 2024 exchange — Correct

Quoted / recovered wording: Recovered wording from prior tally: de-escalation after retaliation.

Quote fidelity: Normalized from prior tally.

What happened later: The later chat tally treated this as a clean hit.

How I resolved it: The reasoning was sharp because it recognized the difference between a controlled punitive exchange and a mutually desired war.

3. Syria's Arab reacceptance would not erase Iran's leverage — Correct

Quoted / recovered wording: Recovered wording from prior tally: normalization optics would not expel Iran from the story.

Quote fidelity: Normalized from prior tally.

What happened later: This direction held up. Diplomatic reintegration did not suddenly neutralize the entrenched Iranian role.

How I resolved it: This is another recurring strength: he resists the fallacy that ceremony equals strategic change.

4. Houthi / Axis attacks would continue as an instrument, not a short-lived spasm — Correct

Quoted / recovered wording: Recovered wording from prior tally: the attack pattern would continue or escalate.

Quote fidelity: Normalized from prior tally.

What happened later: The later tally treated the broad continuation call as correct.

How I resolved it: Good structural reasoning: he understands when a behavior has become strategically embedded.

5. One U.S.–Saudi nuclear or diplomatic packaging call proved too optimistic / too rigid — Incorrect

Quoted / recovered wording: Recovered note from prior tally: one Saudi-related line was scored wrong in the later update.

Quote fidelity: Normalized from prior tally.

What happened later: A specific diplomatic packaging judgment did not resolve the way he implied.

How I resolved it: This suggests his regime-network lens is best on coercion and proxies, less so on transactional diplomatic detail.

6. War risk in the 2025-26 period was greater than many deal-centric readings allowed — Correct

Quoted / recovered wording: Recovered wording fragment from prior tally: the risk of war was greater than the chance of a stable deal while maximalist demands persisted.

Quote fidelity: Fragment recovered from earlier notes.

What happened later: The later tally treated this as one of his stronger calls in the updated round.

How I resolved it: This worked because he prioritized coercive structure over diplomatic theater.

Aggregate diagnosis

Brodsky's sharpest work comes from reading the regime as a coherent pressure machine with multiple adjustable levers: proxies, diplomacy, rhetoric, and procurement.

When he misses, it is often because the conflict lens becomes so dominant that specific diplomatic or transactional details get flattened.

Overall, though, he is a strong medium-horizon analyst of regime behavior.

Mark Dubowitz

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
44 (baseline corpus)	7	4 / 1 / 2	64.3%	Low sample after strict filtering	<i>The updated tally improved his score meaningfully, but the sample remains small and selective. The result says “useful in some coercive-policy calls,” not “highly calibrated forecaster across the board.”</i>

Where the reasoning is strongest

- Good at identifying chronic Iranian coercive behaviors such as hostage-taking, extortion, and sanctions evasion.
- Often right on the limitations of soft diplomatic optimism.

Where the reasoning is weaker

- A tendency toward collapse or catastrophe framing can reduce precision.
- Some of his most memorable lines are rhetorically strong but analytically too binary.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. The 2015 diplomatic process would not end well — Incorrect

Quoted / recovered wording: Recovered wording from prior tally: negotiations were unlikely to end well / pressure would not produce the clean outcome supporters expected.

Quote fidelity: Normalized from prior tally.

What happened later: This was scored against him because the JCPOA was in fact reached in 2015.

How I resolved it: The blind spot was treating his preferred coercive logic as more dominant than the political appetite for a limited deal.

2. Saudi Arabia would step in on oil, though only partially — Partial

Quoted / recovered wording: Recovered note from later tally update: the Saudi compensation call was upgraded to Partial.

Quote fidelity: Normalized from later chat update.

What happened later: Saudi output did rise, but not in a way that cleanly delivered the full strong version of the claim.

How I resolved it: This is a good example of a forecast that captured direction but not magnitude.

3. An Israeli coalition shift and Raisi's ascent in Iran would both occur in 2021 — Correct

Quoted / recovered wording: Recovered note from later tally update: “coalition and Raisi” were counted as correct.

Quote fidelity: Normalized from later chat update.

What happened later: The later tally upgraded this as one of the cleaner additions in the Mark update round.

How I resolved it: This worked because it was a concrete political sequencing call rather than a grand thesis about collapse.

4. Iranian hostage-taking / extortion behavior would continue — Correct

Quoted / recovered wording: Recovered wording from later tally update: hostage-taking and extortion remained a durable pattern.

Quote fidelity: Normalized from later chat update.

What happened later: This was treated as a valid and correct behavior-pattern forecast in the later update round.

How I resolved it: His best calls often look like this: not dramatic singularity, but durable coercive habit.

Aggregate diagnosis

Dubowitz is more useful when he is talking about Iran's persistent coercive toolkit than when he is implying that every diplomatic process is on the verge of a stark binary end-state.

His updated score improved once the tally rewarded durable behavior-pattern calls rather than only dramatic endpoint calls.

Still, the sample is too small to treat the percentage as stable.

Samuel Bendett

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
37 (baseline corpus)	8	4 / 0 / 4	50.0%	Low-to-medium sample, mixed record	<i>A classic mixed-score analyst. He is often early and useful on unmanned-systems diffusion, but novelty bias and a willingness to amplify flashy concepts can cut hard against him.</i>

Where the reasoning is strongest

- Good at seeing diffusion before many mainstream observers do.
- Strong on drone/UxS proliferation and battlefield adaptation.

Where the reasoning is weaker

- Headline-forward amplification can turn concept material into prematurely scoreable “predictions.”
- More vulnerable than the stronger performers here to novelty bias and industrial-promise hype.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. China’s combat-drone presence in the Middle East would grow — Correct

Quoted / recovered wording: Recovered wording fragment: “China’s growing combat drone presence.”

Quote fidelity: Fragment recovered from earlier update notes.

What happened later: The later update treated this as correct, in light of continuing Chinese drone deals and presence.

How I resolved it: This is exactly the kind of diffusion call he does well: a capability moving outward before the security conversation fully catches up.

2. Terrorist drone swarms would emerge soon — Correct

Quoted / recovered wording: Recovered wording fragment: “terrorist drone swarms soon.”

Quote fidelity: Fragment recovered from earlier update notes.

What happened later: This was scored correct in the later update round as militant drone use normalized and scaled.

How I resolved it: Again, this is his best lane: capability diffusion from state to non-state actors.

3. The TALOS / Iron Man-suit line would still matter operationally — Incorrect

Quoted / recovered wording: Recovered wording fragment: “Iron Man suit delivery.”

Quote fidelity: Fragment recovered from earlier update notes.

What happened later: This was later scored wrong once the project’s practical trajectory and cancellation/irrelevance became clear.

How I resolved it: This is the novelty-bias failure mode: concept glamour overpowering procurement reality.

4. A dramatic Russian “death train” type claim would materialize on schedule — Incorrect

Quoted / recovered wording: Recovered wording fragment: “death train not coming in 2019.”

Quote fidelity: Fragment recovered from earlier update notes.

What happened later: This was also later scored wrong.

How I resolved it: The miss reflects a recurring issue in his weaker calls: interesting signal, weak conversion into reliable timeline forecast.

5. Russia would invest seriously in domestic counter-UAS technology — Correct

Quoted / recovered wording: Recovered wording from prior tally: Rostec / domestic C-UAS investment.

Quote fidelity: Normalized from prior tally.

What happened later: That direction was treated as correct in the prior tally.

How I resolved it: This was a sober industrial adaptation call and therefore much stronger than the flashier concept predictions.

6. A quantum-supremacy / leap-type claim from 2020 overstated what actually happened — Incorrect

Quoted / recovered wording: Recovered note from prior tally: a 2020 quantum-supremacy type call was scored wrong.

Quote fidelity: Normalized from prior tally.

What happened later: The later tally marked it down as a miss.

How I resolved it: Same pattern: premature technological maximalism.

Aggregate diagnosis

Bendett is useful when the question is “what capability is diffusing?” rather than “what glossy concept will become real by when?”

His correct calls point to genuine sensitivity to drone and UxS adaptation.

His misses show how quickly tech analysis can drift if product demos, promotional headlines, or eye-catching concepts are treated like mature programs.

Trita Parsi

Base screened corpus	Resolved	C / P / I	Weighted	Confidence	Bottom-line reading
89 (baseline corpus)	26	15 / 2 / 9	61.5%	High sample, mixed-to-moderate accuracy	<i>A large, meaningful sample with a genuinely mixed record. He is often very strong on the limits of pressure and the mechanics of failed coercive diplomacy, but much weaker when the problem turns into hard military escalation or regional security spillover.</i>

Where the reasoning is strongest

- Excellent at seeing when sanctions and maximalist pressure will not deliver the cleaner agreement their advocates promise.
- Good on bargaining structure, diplomatic incentives, and the political economy of failed negotiations.
- Often correct that transactional diplomacy does not automatically produce a stronger replacement arrangement.

Where the reasoning is weaker

- Can underweight the probability of military escalation or regional spillover.
- At times too dismissive of harder-edged security trajectories once deterrence begins to unravel.

Recovered case log

Important: these entries are the recoverable scored predictions from the earlier tally work in this chat. Where exact tweet text is not recoverable from the current visible context, I state that and use normalized prediction wording.

1. 2015 pressure would not produce a magically superior outcome — Correct

Quoted / recovered wording: Recovered wording from prior tally: pressure/leverage arguments against the 2015 talks were overstated.

Quote fidelity: Normalized from prior tally.

What happened later: This part of his worldview held up well: coercive maximalism did not produce a stronger durable deal.

How I resolved it: This is classic Parsi strength—understanding the failure mode of “more pressure = better agreement.”

2. A near-term Israel-Hezbollah war prediction from the mid-2010s did not land — Incorrect

Quoted / recovered wording: Recovered note from prior tally: an Israel-Hezbollah war call was scored wrong.

Quote fidelity: Normalized from prior tally.

What happened later: The later tally counted this as a miss.

How I resolved it: The blind spot here was treating pressure dynamics as more automatically explosive in the short term than they turned out to be.

3. After the U.S. left the JCPOA, pressure still would not yield a better bargain — Correct

Quoted / recovered wording: Recovered wording from prior tally: quitting the deal would not force a stronger replacement on U.S. terms.

Quote fidelity: Normalized from prior tally.

What happened later: That broader thesis held up.

How I resolved it: Again, this is his core competence: bargaining failure under maximalist assumptions.

4. INSTEX would not become a decisive workaround — Correct

Quoted / recovered wording: Recovered note from prior tally: INSTEX was effectively weak / ineffective.

Quote fidelity: Normalized from prior tally.

What happened later: The prior tally treated this as a clean correct call.

How I resolved it: Good institutional realism; he did not overstate a politically popular but thin mechanism.

5. Syria's return to the Arab League would happen despite the broader moral/political objections — Correct

Quoted / recovered wording: Recovered note from prior tally: Syria's return was scored correct.

Quote fidelity: Normalized from prior tally.

What happened later: That call held.

How I resolved it: Strong on regional diplomatic normalization logic.

6. Some “the U.S. will not attack soon” / “no stronger deal is coming” lines were only partly durable once the 2025-26 military track accelerated — Mixed

Quoted / recovered wording: Recovered wording from prior tally: a set of cautious military predictions was mixed or wrong once later conflict outran them.

Quote fidelity: Normalized from prior tally.

What happened later: His skepticism about a better deal remained sensible, but the harder military path was underweighted in parts of the record.

How I resolved it: This is the central tradeoff in his work: very good on failed diplomacy, less strong on deterrence breakdown.

7. Russia-Iran cooperation would deepen — Correct

Quoted / recovered wording: Recovered note from prior tally: Russian-Iranian cooperation deepened and was counted correct.

Quote fidelity: Normalized from prior tally.

What happened later: This direction held.

How I resolved it: A good call when he stayed at the level of strategic alignment rather than narrow tactical prediction.

Aggregate diagnosis

Parsi’s record is best read as a split profile: strong on the diplomatic mechanics of failure, weaker on the military consequences that can follow when diplomacy fails for too long.

That does not make him a weak analyst; it makes him domain-skewed.

The size of his sample makes that conclusion sturdier than the topline for the lower-volume accounts.

Tal Inbar

Baseline screened corpus	Resolved (updated)	C / P / I	Weighted	Confidence	Bottom-line reading
23 baseline-screened tweets + later raw addendum	23	14 / 2 / 7	65.2%	Medium	<i>Technically sharp on launch cadence, site readiness, and propulsion conservatism; far weaker on spectacular escalation or deadline-driven brinkmanship.</i>

Integration note

The earlier cross-expert tally in this chat treated Tal separately and produced a 7-call resolved ledger. The later raw-prediction audit added 16 more standalone scoreables and explicitly said it was a standalone increment rather than a merged master row.

This chapter therefore gives the merged updated score for Tal at the top, while reproducing in detail the 16 recovered addendum cases that are available in quote-level form from the current context. The earlier 7 baseline cases are already absorbed into the topline score but are not all reproducible quote-for-quote from the current visible context.

Aggregate reasoning profile

Tal Inbar's work is at its best when the artifact itself constrains the analysis: the pad, the tower, the trajectory, the propulsion choice, the cadence after failure, or the political calendar of a prestige launch attempt.

His strongest contrarian wins come from preferring base rates and physical evidence over the excitement cycle of social media. His misses cluster elsewhere: once the story becomes about a Pacific atmospheric test, a dramatic Christmas surprise, or an Iranian nuclear-threshold retaliation, the analysis becomes more imaginative and less disciplined.

The deepest blind spot is escalation bias tied to authoritarian signaling. He often correctly sees that prestige matters to Pyongyang or Tehran, but then sometimes over-infers that prestige pressure will force the most dramatic available move, even when technical readiness or strategic caution argue otherwise.

Recovered quote-level addendum cases

TI-01. The burst of NOTAM-related speculation around Iran's new launch complex in January 2015 was unlikely to produce an actual launch in the near term. — True

Tweet quote: “Many NOTAM's have been issued during the last years, but ended in nothing. So I would advise- not to hold our breath...”

Later record: The feared/implied launch wave did not materialize in early 2015. The Imam Khomeini Space Center only emerged as an operational launch site later, with the complex officially reopened and associated with a Simorgh launch in July 2017.

Resolution analysis: This was a classic base-rate call. Inbar was pushing back against OSINT excitement by using two simple filters: (1) Iran had a history of NOTAMs that did not culminate in launches, and (2) the physical readiness of the site still looked doubtful. The reasoning was strong because it preferred observed program behavior over hype. He did not say 'never'; he said 'not now'—and that timing judgment held up.

TI-02. Following failure, Kim Jong Un would intensify attempts to produce a visible missile success rather than pause quietly. — True

Tweet quote: “#NORTHKOREA Missile failure will hurt KJU credibility as a threat, and increase his effort to show success.”

Later record: North Korea did return rapidly to testing. On 22 June 2016, after a string of Musudan failures, it conducted new launches, including one that was widely treated as a partial success.

Resolution analysis: This call was not about hardware identification; it was about the political economy of a prestige weapons program. Inbar correctly recognized that repeated public failure raises the regime's incentive to retry quickly and visibly. The logic was sharp because it linked engineering cycles to authoritarian prestige management.

TI-03. Another North Korean missile attempt would follow quickly after the 15 October 2016 failure. — True

Tweet quote: “#NorthKorea launched a ballistic missile. US STRATCOM reports it was a failure. I expect to see another try Soon.”

Later record: A further North Korean missile launch followed only days later, on 20 October 2016, and it too failed.

Resolution analysis: Again, the strength here was rhythm-reading. Inbar understood that a failure in a politically important campaign does not usually end the campaign; it compresses the timeline to the next attempt. This is a useful pattern-recognition skill in missile analysis.

TI-04. North Korea would soon test the Hwasong-13 / KN-08-class ICBM as a dramatic answer to recent failures. — False

Tweet quote: “After 2 failures, #NorthKorea must pull a rabbit out of the hat. I'm betting on an HS-13 test soon.”

Later record: North Korea did not move to a KN-08/Hwasong-13 flight-test sequence. Open-source tracking later emphasized that the Hwasong-13 family was never flight-tested.

Resolution analysis: This was the same prestige logic as earlier—but now overextended. The blind spot was assuming that political need for a dramatic success would override technical readiness. That is often wrong in missile programs. A regime may want a spectacular new system, but if the system is not ready, desire does not create hardware. The call seems to have been pulled toward the most exciting possibility rather than the most mature one.

TI-05. North Korea's first true ICBM flight would more likely be lofted eastward from an operational missile area than flown southward from the Sohae space center. — True

Tweet quote: “Future NK launch of an ICBM probably be a lofted trajectory to the East, rather than a launch from Sohae space center to the South.”

Later record: North Korea's July 2017 Hwasong-14 launches were lofted flights toward the East Sea/Sea of Japan, not southbound launches from Sohae.

Resolution analysis: This was one of Inbar's best calls in the batch. He integrated test-geometry logic, political risk management, and the difference between a space-launch site and an operational missile campaign. It was also mildly contrarian, because public attention often overfixates on Sohae. He anticipated the actual test logic instead of the visually familiar location.

TI-06. If the United States did not respond in a materially stronger way to the 29 August 2017 overflight of Japan, North Korea would repeat the same basic pattern soon. — True

Tweet quote: “NK missile flight over Japan is an Archimedean point for the US. No REAL ACTION - and we will see more missiles on the same path soon.”

Later record: On 15 September 2017, North Korea launched another Hwasong-12 that again flew over Japan.

Resolution analysis: The reasoning here was sharp because it treated missile launches as signaling sequences, not isolated events. Inbar inferred that the regime had discovered a coercive act it could repeat at manageable cost. That is a strong operational reading of adversary behavior.

TI-07. An atmospheric North Korean nuclear test over the Pacific was moving from fringe speculation to likely future reality. — False

Tweet quote: “NK nuclear test over the Pacific: First a wild idea, then came articles and analysis. Soon it won't be IF but WHEN.”

Later record: No such Pacific atmospheric nuclear test occurred. North Korea's latest confirmed nuclear test remains the 2017 underground test.

Resolution analysis: This miss captures a recurring weakness in the batch: once the discussion shifted from engineering and test geometry to spectacular escalation, Inbar's imagination ran ahead of the threshold the regime was actually willing to cross. The blind spot was underestimating the extraordinary diplomatic, military, and environmental costs of a Pacific atmospheric nuclear test.

TI-08. The long-duration November 2017 launch was probably another Hwasong-14 rather than a genuinely new ICBM. — False

Tweet quote: “Japanese sources said NK missile flew for 50 minutes (!). If true, it is probably HS-14.”

Later record: The launch turned out to be the first Hwasong-15 flight, not another Hwasong-14.

Resolution analysis: This was a sparse-data classification error. The problem was not irrationality; it was that Inbar used a conservative identification heuristic at the exact moment when the program had actually jumped to a larger system. In other words, the same conservatism that protected him from hype in other cases briefly made him too slow to credit the step-change.

TI-09. The Hwasong-15 was unlikely to represent a North Korean move to solid-fuel ICBMs. — True

Tweet quote: “Not likely that HS-15 is Solid propelled. siin we will Know.”

Later record: The Hwasong-15 was understood as a liquid-fuel ICBM, not a solid-fuel system.

Resolution analysis: This is another example of Inbar’s engineering conservatism at its best. In periods of rapid advance, analysts often over-project the next breakthrough. He resisted that temptation. That restraint was valuable because solid-fuel ICBMs would have implied a much larger leap in survivability and readiness.

TI-10. North Korea would soon end its 2018 testing pause and return to missile launches as diplomacy soured. — False

Tweet quote: “One (maybe two) more threat by Trump towards Kim Jong Un and we might see the FIRST missile test by NK for 2018.” || “OK. The clock has started for NK first missile test for 2018. Back to business as usual.”

Later record: North Korea did not conduct a ballistic-missile test in the remainder of 2018. The test pause lasted until 2019.

Resolution analysis: This was a classic 'reversion to normal' mistake. Inbar assumed the structural logic of missile politics would quickly reassert itself. What he underweighted was the usefulness of the moratorium itself as a negotiating instrument. Kim did not need a launch immediately; the pause was itself part of the strategy.

TI-11. The Sohae reconstruction scare of March 2019 did not by itself indicate an imminent ICBM or space-launch resumption. — True

Tweet quote: “Dismantling some light structures at NK's Sohae Satellite Launching Station was NOT a step towards ballistic missiles disarmament exactly as re-assembling the structures isn't an oment sign of resuming ICBM tests.” || “As for now, the whole thing looks like ‘Much Ado About Nothing’.”

Later record: Despite the alarm, there was no immediate Sohae-based ICBM or SLV launch. Contemporary official reporting also said there was no sign of an imminent launch.

Resolution analysis: This was a very good OSINT discipline call. Inbar distinguished between visible site activity and actual launch readiness, and between symbolic compliance gestures and operational decisions. That separation is exactly where a lot of commentary goes wrong.

TI-12. North Korea's late-2019 'Christmas gift' rhetoric would likely culminate in a significant pre-Christmas test or surprise. — False

Tweet quote: “10 days till Christmas. Time for 1-2 additional NK tests before the much anticipated move.” || “A real Christmas SURPRISE might come before Christmas. Just saying...”

Later record: There was no dramatic Christmas missile surprise in 2019.

Resolution analysis: This miss looks a lot like TI-10 and TI-07: deadline rhetoric and strategic theater pulled the analysis toward a dramatic event. The blind spot was assuming that public suspense must resolve in public action. Often, especially in coercive diplomacy, suspense is itself the instrument.

TI-13. Iran would try to launch Zafar during the Ten Days of Dawn window in early February 2020. — True

Tweet quote: “haunted by the unreliable Simorgh SLV, Iran built TWO identical ZAFAR satellites, to be launched 'SOON', AKA during the first 10 days of February.”

Later record: Iran launched Zafar on 9 February 2020, squarely inside the predicted window.

Resolution analysis: This is a good example of Inbar's feel for Iranian political choreography. The reasoning was not purely technical; it was institutional. He recognized that symbolic national anniversaries structure when Tehran prefers to unveil or attempt prestige space activities.

TI-14. Zuljanah would be an all-solid, TEL-based IRGC launcher developed mainly as an ICBM in satellite-launch clothing, and it would operate from Shahroud. — Mixed

Tweet quote: “An Iranian three stage all solid propelled SLV on TEL, developed by IRGC will be ALSO a satellite launcher, but mainly an ICBM. Same dimentions as Simorgh - Zoljanah will be impressive, and Shahrud site will be the launch site.”

Later record: Open reporting later described Zuljanah as a mixed-propulsion launcher rather than an all-solid one. Open-source reporting on site association has also been more mixed than the tweet's certainty implied. At the same time, his larger concern—that Iranian military and space-launch work sit very close to ballistic-missile development—was not frivolous and remains a central analytical issue.

Resolution analysis: The broad instinct—strong military-space dual use—was directionally serious and not frivolous. But the specific package was too exact and overconfident. Inbar folded fuel type, institutional ownership, intended role, and launch site into one strong claim. Once several sub-claims are bundled together, the chance that one or more will fail rises sharply.

TI-15. North Korea would field a new satellite-launch vehicle in 2023. — True

Tweet quote: “Still to come - some SLBM TESTS (new systems) and drills. Also we can expect a new SLV. 2023 is an interesting year...”

Later record: North Korea indeed introduced the Chollima-1 launcher in 2023, with failed attempts in May and August and a successful orbital launch in November.

Resolution analysis: This worked because it was grounded in program direction rather than headline drama. North Korea had publicly elevated reconnaissance-satellite ambitions, so a new launcher campaign was a plausible institutional next step.

TI-16. Iran might retaliate for Fakhrizadeh's assassination with a much more dramatic step—up to a demonstrated nuclear weapon. — False

Tweet quote: “Iran's 'severe revenge' over Fakhrizadeh assassination could be anything from counter-US base or proxy operation against Israel. But it could be MUCH MORE DRAMATIC and could lead to immunity and survivability of the regime: Demonstration of a nuclear weapon - KJU style.”

Later record: No Iranian nuclear demonstration followed. Even much later, U.S. intelligence reporting still held that Iran had not reauthorized a nuclear-weapons program.

Resolution analysis: This is the clearest example in the batch of escalation imagination outrunning observed thresholds. The flaw was treating the political need for revenge as if it were enough to force a nuclear threshold crossing. But a nuclear demonstration is not just a retaliation option—it is a strategic regime decision with enormous costs and technical requirements.

What the merged Tal score means

The merged Tal score should not be read as an average analyst score in the abstract. It is a composite of two very different domains inside one feed: technically grounded missile/space OSINT, and much shakier threshold-forecasting about dramatic state behavior.

A user of Tal's work should trust him more when the question is “Is this site actually ready?”, “What propulsion category is this likely to be?”, “Will the regime retry after an embarrassing failure?”, or “Is OSINT overreading a picture?”

A user should trust him much less when the question becomes “What spectacular threshold act will Pyongyang or Tehran choose next simply because it would be dramatic?”

John J. Mearsheimer

Resolved: 15 | Correct: 10 | Partial: 3 | Incorrect: 2 | Weighted hit rate: 76.7% | Confidence note: Moderate-high confidence; many entries in the source file were bibliographic listings, paywalled stubs, or fetch errors, so only text-bearing recoverable calls were scored.

Where the reasoning is strongest

- Excellent at identifying when a state sees an issue as a core security interest and therefore will absorb pain rather than back down.
- Very strong on occupation logic: he distinguishes winning the initial war from controlling the political order that follows.
- Strongest when he reasons from structure — geography, power asymmetry, nationalism, deterrence, and great-power rivalry — rather than from rhetoric.

Where the reasoning is weaker

- Sometimes over-specifies how a structurally sound argument will cash out in a discrete endpoint or timeline.
- Can underweight the ability of institutions, frozen conflicts, or outside guarantors to preserve an ugly but durable equilibrium.
- When the broad structural diagnosis is right, he can still be too categorical about second-order consequences, especially on proliferation and balancing coalitions.

Recovered case log

Important: unlike several earlier chapters in this report, the Mearsheimer section is built from recoverable article text in the uploaded JSON rather than from prior tally notes alone. Where possible, I therefore quote exact wording from the article text. When a forecast had to be normalized from the extraction ledger rather than quoted from a single clean sentence, I say so explicitly.

1. 1991 Gulf War: Kuwait could be liberated in less than a week with fewer than 1,000 U.S. fatalities — Correct

Quoted / recovered wording: "The U.S. military can liberate Kuwait in less than a week and suffer relatively few casualties -- probably less than 1,000 fatalities."

Quote fidelity: Exact excerpt from the recovered New York Times text.

What happened later: The coalition ground campaign lasted about 100 hours, and official U.S. battle deaths were 148. This was a clean hit on both duration and order-of-magnitude casualties.

How I resolved it: This call worked because he was not mesmerized by Iraqi fortifications. He focused on the mismatch in mobile armored warfare, training, night-fighting capability, artillery

quality, and coalition air dominance. The sharp part of the reasoning was distinguishing static defensive appearance from operational inferiority in maneuver war.

2. 1993 Ukraine: Ukraine was unlikely to transfer its nuclear weapons to Russia — Incorrect

Quoted / recovered wording: "Second, it is unlikely that Ukraine will transfer its remaining nuclear weapons to Russia, the state it fears most."

Quote fidelity: Exact excerpt from the recovered Foreign Affairs text.

What happened later: Ukraine did transfer the inherited warheads and became non-nuclear by 1996. The strategic logic he saw was real, but the forecast about what Kyiv would actually do was wrong.

How I resolved it: The blind spot was political choice under pressure. He correctly identified the security logic favoring retention, but he underweighted elite preferences for international integration, the pressure applied by Washington and Moscow, and the willingness of Ukrainian leaders to trade deterrent assets for assurances and diplomatic positioning.

3. 1993 Ukraine: a non-nuclear Ukraine would be exposed to Russian coercion and possible war — Correct

Quoted / recovered wording: "That means ensuring that the Russians ... do not move to reconquer it. Ukraine cannot defend itself against a nuclear-armed Russia with conventional weapons ... Ukrainian nuclear weapons are the only reliable deterrent to Russian aggression."

Quote fidelity: Composite of adjacent exact excerpts from the recovered Foreign Affairs text.

What happened later: Russia seized Crimea in 2014, supported armed conflict in the Donbas, and launched a full-scale invasion in 2022. The specific war path took longer than he implied, but the core vulnerability diagnosis was strongly vindicated.

How I resolved it: This is one of his signature successes. He treated formal assurances and legal language as weak substitutes for hard deterrent capability. That was contrarian in the optimistic 1990s, but it proved sharp because it started from power asymmetry, geography, and the possibility that Russia would again define Ukraine as a vital sphere-of-influence issue.

4. 1997 Bosnia: absent partition, war would erupt again after U.S. withdrawal — Incorrect

Quoted / recovered wording: "But war will erupt again when we leave ... Only a managed partition can let the United States leave without triggering a new war."

Quote fidelity: Exact excerpt from the recovered New York Times text.

What happened later: Bosnia remained fragmented, ethnically divided, and politically dysfunctional, but large-scale war did not resume after Dayton. The conflict was frozen rather than reignited.

How I resolved it: The reasoning saw the hatred and the institutional weakness clearly, but it overestimated how inevitably those conditions would revert to war. The blind spot was the staying power of external supervision, deterrence, and exhaustion. He was right that Dayton built an unstable order; he was wrong that instability had to cash out as renewed warfare on the timeline he expected.

5. 1998 Kosovo: the Holbrooke deal would fail sooner rather than later — Correct

Quoted / recovered wording: "The deal arranged by Richard Holbrooke is likely to fail sooner rather than later."

Quote fidelity: Exact excerpt from the recovered New York Times text.

What happened later: The cease-fire diplomacy collapsed, Serbian operations resumed, and NATO launched its air campaign in 1999. The core forecast was right.

How I resolved it: This was sharp because he identified the incompatibility built into the deal. Albanians wanted independence, not autonomy inside Serbia; Serbs feared that autonomy was merely a way-station to losing Kosovo altogether. When each side reads the same arrangement as a mortal long-term threat, compliance is fragile from the outset.

6. 1998 South Asia: Pakistan would follow India down the nuclear path — Correct

Quoted / recovered wording: "it is surely in India's self-interest to build a nuclear deterrent -- as it will be for Pakistan to follow suit."

Quote fidelity: Exact excerpt from the recovered New York Times text.

What happened later: Pakistan tested nuclear weapons later in May 1998, just after India. This was a straightforward hit.

How I resolved it: The reasoning was simple but powerful. He treated Indian and Pakistani nuclear choices as interactive rather than isolated. Once India crossed the threshold publicly, Pakistan faced an acute deterrence and prestige problem. This was an excellent example of second-move logic.

7. 2001 Afghanistan: a U.S. ground war would generate insurgency and prove unwinnable at acceptable cost — Correct

Quoted / recovered wording: "The Taliban would launch a guerrilla resistance from the countryside. It is unlikely that the United States could win this armed struggle at any reasonable cost."

Quote fidelity: Exact excerpt from the recovered New York Times text.

What happened later: The war became a long counterinsurgency, the Taliban survived, and after twenty years the United States withdrew while the Taliban returned to power. This was one of his strongest calls.

How I resolved it: He saw the occupation problem very early. The sharp insight was that destroying the incumbent regime is not the same thing as solving the political problem. He also understood how terrain, sanctuary, nationalism, and outside support interact to keep insurgencies alive long after conventional defeat.

8. 2001 Israel/Palestine: Israel would not make the concessions needed for a viable Palestinian state anytime soon — Correct

Quoted / recovered wording: "It is hard to imagine the Palestinians accepting such a state ... Israel won't agree to such measures anytime soon."

Quote fidelity: Composite of exact excerpts from the recovered New York Times text.

What happened later: No final-status agreement created a viable sovereign Palestinian state, and the conflict remained unresolved. The pessimism proved well grounded.

How I resolved it: This was not just cynicism. He traced the incompatibility between a truly viable Palestinian state and Israel's perceived security requirements. The logic was sharp because it treated the dispute as a collision between sovereignty and security, not as a problem of atmospherics or negotiator chemistry.

9. 2003 Iraq: overthrowing Saddam would require years of occupation and cost far more than advertised — Correct

Quoted / recovered wording: "toppling Saddam Hussein would cost at least \$50 billion to \$100 billion ... the actual cost of this war would most likely be much larger."

Quote fidelity: Composite of exact excerpts from the recovered New York Times text.

What happened later: The United States occupied Iraq for years; the war and its downstream costs ran vastly beyond the original public estimates. On the cost and duration logic, this was a strong hit.

How I resolved it: The key reasoning strength was separating regime collapse from postwar order. He anticipated that occupation, state rebuilding, insurgency, and regional blowback would dominate the cost curve. Many advocates reasoned from the short war to the cheap war; he did not.

10. 2005 China: China's rise would not be peaceful and would trigger intense U.S.-China security competition — Partial

Quoted / recovered wording: "If China continues its impressive economic growth over the next few decades, the US and China are likely to engage in an intense security competition ... Most of China's neighbours ... will join with the US to contain China's power."

Quote fidelity: Composite of exact excerpts from the recovered Australian op-ed text.

What happened later: The first half of the forecast was strongly vindicated: U.S.-China rivalry became central to world politics. The coalition claim was more mixed. Many states deepened security ties with Washington, but the predicted balancing front has been uneven, hedged, and less unified than his strongest formulation implied.

How I resolved it: This is a classic partial. He nailed the structural rivalry mechanism — rising power, incumbent hegemon, regional ambitions, and fear among neighbors. The overstatement lay in assuming that balancing behavior would cohere more cleanly and more uniformly than it has under conditions of economic interdependence and strategic hedging.

11. 2009 Afghanistan: Obama would escalate and the United States would be in Afghanistan for the long haul — Correct

Quoted / recovered wording: "The United States is in Afghanistan for the long haul."

Quote fidelity: Exact excerpt from the recovered Foreign Policy text, reinforced by the companion Newsweek argument that the war had become a quagmire.

What happened later: Obama approved the surge in 2009, and the war continued until 2021. This was a clear hit on both direction and staying power.

How I resolved it: He combined strategic logic with domestic political logic. The sharp part was recognizing that presidents often deepen failing wars because the political cost of visibly "losing" seems worse in the short run than prolonging the conflict. That lock-in mechanism mattered enormously.

12. 2014 Ukraine: Western pressure would worsen the crisis; Russia could seize territory and would not back down — Correct

Quoted / recovered wording: "Instead of resolving the dispute, it will lead to more trouble ... It could invade eastern Ukraine or annex Crimea ... Mr. Putin is unlikely to back down."

Quote fidelity: Composite of exact excerpts from the recovered New York Times text.

What happened later: Russia annexed Crimea in 2014, supported separatist war in eastern Ukraine, and later escalated further. This was one of the most famous successful Mearsheimer calls.

How I resolved it: Again, the sharpness came from structural realism. He treated NATO expansion and Western alignment on Russia's border as a problem Moscow would interpret through hard-security logic. Many analysts moralized Putin; he modeled what Russia would do when faced with what it saw as encroachment in a vital zone.

13. 2015 Ukraine: arming Ukraine would escalate the war and Russia would counter-escalate rather than cave — Partial

Quoted / recovered wording: "Sending weapons to Ukraine will not rescue its army and will instead lead to an escalation in the fighting ... Russia would counter escalate ... This coercive strategy is also unlikely to work, no matter how much punishment the West inflicts."

Quote fidelity: Composite of exact excerpts from the recovered New York Times text.

What happened later: Escalation certainly occurred, and Russia did not cave under pressure. But the stronger implication that Western military aid could not materially help Ukraine was too categorical, especially after 2022 when large-scale aid clearly mattered for Ukrainian survival and battlefield effectiveness.

How I resolved it: He got the Russian resolve and escalation logic right, but the blind spot was treating aid mainly as futile coercion rather than as a tool that could substantially improve Ukraine's defensive capacity even without forcing outright Russian capitulation.

14. 2019 Iran: maximum pressure would not produce capitulation and would instead generate retaliation and escalation pressure toward the bomb — Partial

Quoted / recovered wording: "There is no evidence that Iran is likely to capitulate ... The Iranians are likely to launch more covert attacks ... Iran will double down, which almost certainly means it will move to build its own nuclear arsenal."

Quote fidelity: Composite of exact excerpts from the recovered New York Times text.

What happened later: The non-capitulation and retaliation parts were strongly borne out. Iran and its network retaliated in the Gulf theater and the spiral intensified. The stronger claim that Tehran would move all the way to an actual nuclear arsenal remains unproven on the available public evidence.

How I resolved it: This was another structure-right, endpoint-too-strong case. He correctly understood nationalism, coercive failure, and escalatory spirals. The miss was in moving from "pressure increases incentives for deterrence" to a near-certain prediction of weaponization.

15. 2022 Ukraine: there was little chance of meaningful compromise and the war would settle into a prolonged stalemate with grave escalation risks — Correct

Quoted / recovered wording: "there is little chance of a meaningful compromise" and observers "foresee a bloody stalemate. They are right about that."

Quote fidelity: Exact excerpts from the recovered Foreign Affairs text.

What happened later: As of April 2026 there is still no durable settlement, the conflict remains protracted, and escalation risks — though managed — never disappeared. This was a strong hit.

How I resolved it: The reasoning was sharp because it looked at war aims rather than press-release diplomacy. When both sides have maximalist objectives and fear the consequences of losing, compromise space collapses. He saw that the war was structurally sticky long before optimism about near-term bargaining fully faded.

Aggregate diagnosis

Mearsheimer's record in this ledger is strongest where structural realism has the most predictive leverage: great powers reacting to threats near their borders, occupations generating insurgencies, rivals balancing against rising power, and states refusing to capitulate when they define the stakes as existential.

What makes his best calls impressive is not that they are flashy. It is that he repeatedly strips away moral theater, diplomatic branding, and wishful thinking and asks a harsher question: what will states do when survival, deterrence, prestige, and regional power are on the line? That is why his Afghanistan, Iraq, and Ukraine warnings aged so well.

His main failure mode is over-certainty in the final step. Sometimes he moves too quickly from a sound structural diagnosis to a specific endpoint: Ukraine would probably keep its nuclear weapons, Bosnia would probably return to war, Iran would probably move all the way to an actual bomb, or Asia would cohere into a more unified balancing coalition than it in fact has. In

short, his macro logic is often excellent; his discrete endpoint calibration is occasionally too sharp-edged.

Cross-expert synthesis

The strongest performers in this chat are not the same kind of analysts. Danny Citrinowicz is the most robust all-around forecaster in the current ledger because he combines a large resolved sample with consistently strong judgment about regional escalation pathways. Behnam Ben Taleblu has the highest hit rate, but on a much smaller sample. Alex Plitsas is sharp on directional conflict logic. Jason Brodsky is strong on regime-network behavior and sustained proxy pressure.

The mixed middle of the table is also informative. John J. Mearsheimer now lands in the strong upper-middle: very good on long-horizon structural calls about occupation, deterrence, spheres of influence, and great-power rivalry, but sometimes too categorical on the exact endpoint. Tal Inbar is highly usable in the technical-operational lane but much less so in deadline- or spectacle-driven strategic forecasting. Mark Dubowitz improves once the ledger rewards durable coercive-behavior calls rather than only dramatic endpoint claims. Trita Parsi's record is large enough to show a real pattern: good on the limits of pressure and diplomacy, weaker on military spillover. Samuel Bendett is a textbook diffusion analyst with a novelty-bias problem.

In practical use, the best way to exploit this group is not to ask “who is best overall?” but “which analyst is best for which class of uncertainty?” For coercive diplomacy and sanctions adaptation, Ben Taleblu and Parsi are useful in very different ways. For regional war trajectories, Citrinowicz and Brodsky are stronger. For technical missile/space OSINT, Tal is valuable—but only once his escalation-imagination is discounted. For long-cycle structural-realist warnings about occupations, great-power red lines, and escalation ladders, Mearsheimer is highly usable.

Recommendation on future updates

For future batches, the cleanest workflow is a master ledger with one row per resolved scoreable prediction, one expert field, one status field (Correct / Partial / Incorrect), and one normalization field. That prevents score inflation from duplicates and makes later master-report generation easy.

If the raw tweet dumps are preserved alongside the master ledger, the quote-level self-contained report can later be made fully exhaustive rather than partially reconstructed from tally notes.